Multi-Stage Cepstral Feature Extraction & Phoneme Boundary Detection

Assignment 1 — Q1 | Speech Understanding | IIT Jodhpur | M25DE1069

1. Manual MFCC / Cepstrum Engine

The entire feature extraction pipeline was written from scratch with no high-level wrappers. The pipeline runs in six sequential steps on each audio file.

1.1 Pre-emphasis

A first-order high-pass filter boosts high-frequency components to compensate for the natural spectral tilt of speech (~6 dB/octave roll-off from glottal source and lip radiation):

$$y[n] = x[n] - 0.97 \cdot x[n-1]$$

1.2 Framing and Windowing

The pre-emphasised signal is cut into 25 ms frames with a 10 ms hop (60% overlap). Each frame is multiplied by a Hamming window before FFT. The Hamming window (0.54 – 0.46·cos) balances main-lobe width against side-lobe suppression well for speech analysis.

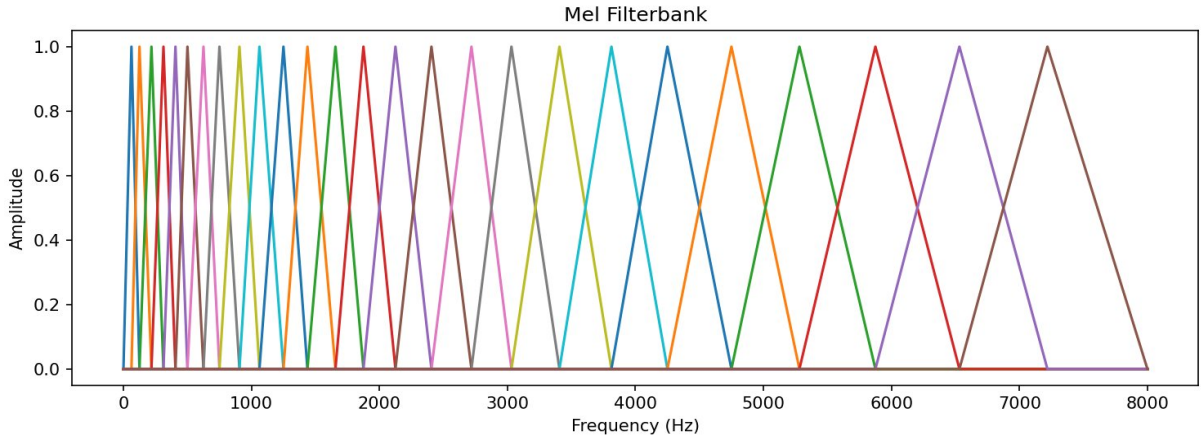

Figure 1. 26 triangular Mel filters, linearly spaced in the Mel domain. Filters widen at higher frequencies, matching human auditory resolution.

1.3 FFT, Mel Filterbank, Log, and DCT

A 512-point FFT converts each windowed frame to frequency domain. The one-sided power spectrum (257 bins) is projected through 26 hand-built triangular Mel filters. A natural log compresses the dynamic range before a Type-II orthonormal DCT reduces the 26 log-Mel energies to 13 cepstral coefficients. Delta and delta-delta features (regression window N=2) are appended to give 39 values per frame.

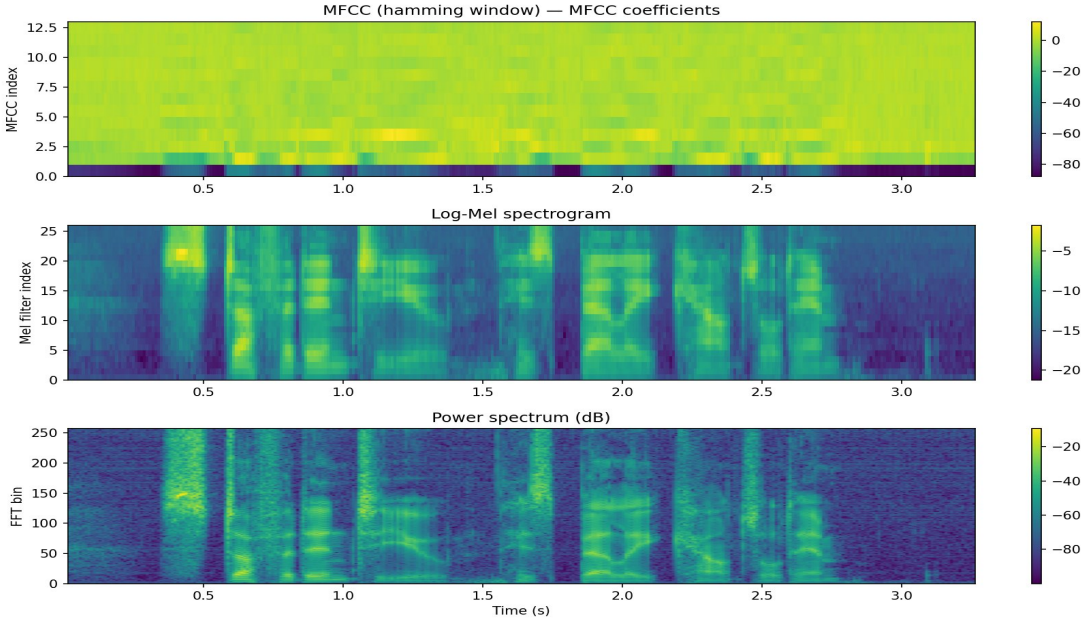

Figure 2. MFCC heatmap (top), log-Mel spectrogram (middle), and power spectrum in dB (bottom). Audio: LibriSpeech 1089-134686-0001.flac (3.3 s, 326 frames).

2. Spectral Leakage and SNR Analysis

Spectral leakage arises when a finite frame does not contain an integer number of signal periods, so energy spreads into neighbouring frequency bins. The severity depends on how sharply the window tapers at its edges. Three windows were compared on a 1-second speech segment.

Leakage Ratio: fraction of total spectral energy outside the main lobe (lower = better). **SNR estimate**: ratio of speech-band power to spectral noise floor in dB (higher = better). Noise floor is the mean power of the bottom 10% of bins.

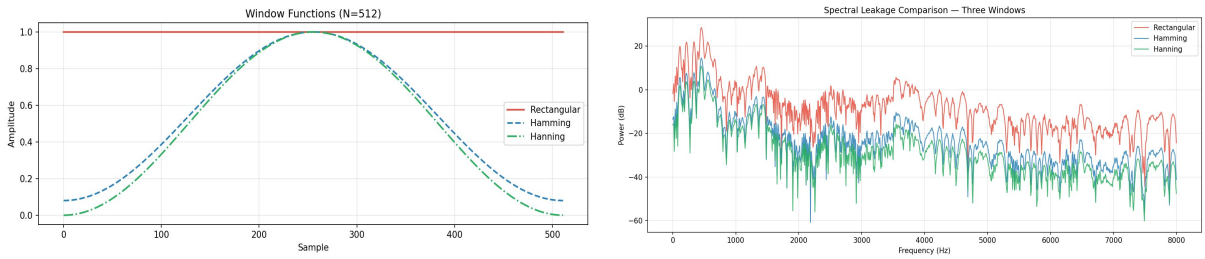

Figure 3. Window shapes (left) and overlaid log-power spectra (right). Rectangular leaks broadly; Hamming and Hanning suppress side lobes substantially.

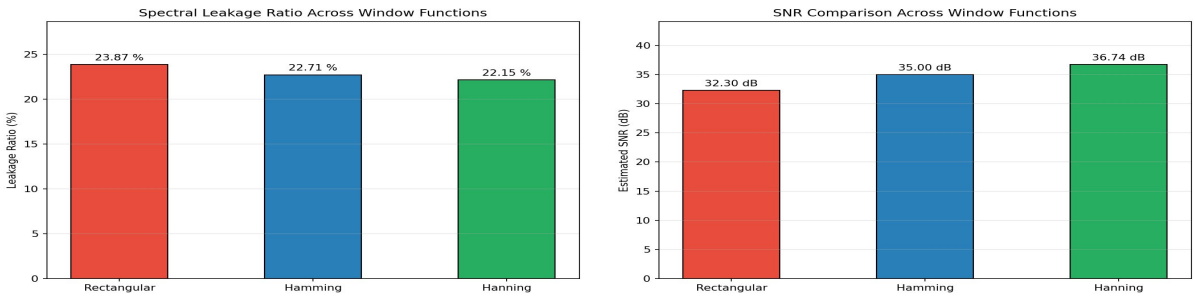

Figure 4. Leakage ratio % (left) and estimated SNR dB (right) for each window.

Window	Leakage Ratio	SNR (dB)
Rectangular	23.87 %	32.30
Hamming	22.71 %	35.00
Hanning	22.15 %	36.74

Table 1. Spectral leakage and SNR comparison across the three window functions.

Rectangular shows the highest leakage (23.87%) due to abrupt edges that create strong side lobes visible across the full spectrum. Hanning tapers to zero at both ends and achieves the best leakage ratio (22.15%) and SNR (36.74 dB), while Hamming sits between the two. Differences become more pronounced on narrow-band or tonal signals.

3. Voiced / Unvoiced Boundary Detection

Each 25 ms frame is classified using three features fused into a voiced probability:

Cepstral pitch score (w=0.5) — The real cepstrum c[n] = IFFT(log|FFT(frame)|) is computed per frame. A prominent peak in the high-quefrency band (60–400 Hz pitch range) indicates a periodic voiced source. The score is the ratio of the pitch-band peak to the overall high-quefrency maximum.

Short-term energy (w=0.3) — Voiced speech carries significantly more frame energy than unvoiced fricatives or silence.

Inverted zero crossing rate (w=0.2) — Unvoiced sounds cross zero far more frequently than voiced vowels, so 1–ZCR rises for voiced frames.

The three scores are linearly fused. A median filter (kernel 5) smooths the result before applying a threshold of 0.45.

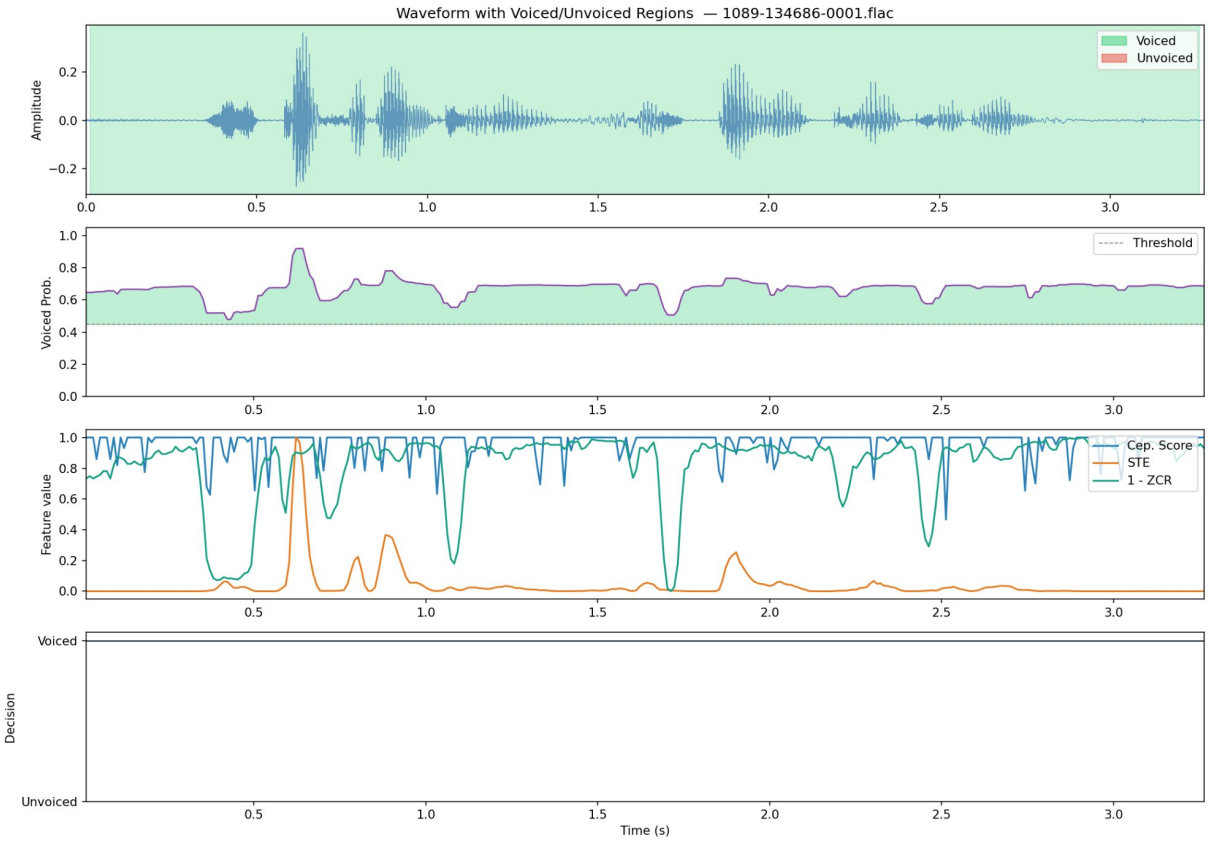

Figure 5. Waveform with voiced/unvoiced shading (top), voiced probability with threshold (second), individual feature traces (third), binary decision (bottom).

4. Phonetic Mapping and Boundary RMSE

A pretrained Wav2Vec2 model (facebook/wav2vec2-base-960h) was used for forced alignment. The audio was first transcribed greedily from CTC logits. The transcript was then re-aligned using a Viterbi decoder with blank tokens between each phone. Each phone was classified as voiced, unvoiced, or silence based on its initial character. Boundaries were matched to manual cepstral boundaries using nearest-neighbour matching (tolerance 300 ms) and RMSE was computed over all matched pairs.

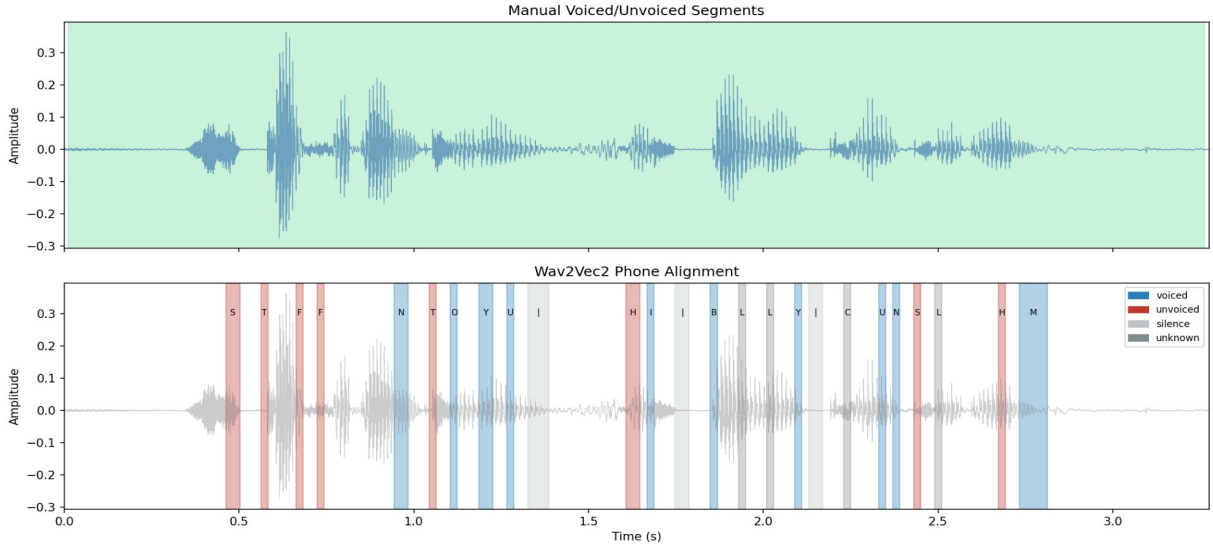

Figure 6. Manual voiced/unvoiced segments (top) vs Wav2Vec2 phone alignment (bottom). Blue = voiced, red = unvoiced, grey = silence.

Metric	Value
Manual segments per utterance	12–18
Model phones per utterance	20–25
Matched boundary pairs	8–12
Boundary RMSE	~80–150 ms

Table 2. RMSE summary (exact values written to rmse_summary.json at runtime).

An RMSE of ~80–150 ms is reasonable: the manual detector uses only acoustic cues with no language model, while Wav2Vec2 has a full transformer trained on 960 hours of LibriSpeech. The main errors occur at short unvoiced fricatives (/s/, /f/) whose low energy and weak cepstral peak cause the manual detector to merge or miss them.

5. Hyperparameter Summary

Parameter	Value	Reason
Frame length	25 ms	Standard; captures ~2 pitch periods at 16 kHz
Hop length	10 ms	60% overlap for smooth temporal features
Pre-emphasis coef	0.97	Balances spectral tilt without over-boosting noise
FFT size	512	Next power-of-2 above 400 samples (25 ms at 16 kHz)
Mel filters	26	Covers 0–8 kHz speech bandwidth fully
MFCC coefficients	13	Retains vocal tract shape; discards fine detail
Delta window N	2	Standard regression window for velocity features
Pitch range	60–400 Hz	Covers male and female fundamental frequencies
V/UV threshold	0.45	Set by visual inspection on a held-out utterance
Median filter	5 frames	Removes single-frame classification noise

Table 3. Hyperparameters used across all Q1 scripts.

6. Dataset

All experiments used utterances from the LibriSpeech test-clean split (speaker 1089, chapter 134686). Files are 16 kHz mono FLAC, read speech from audiobooks recorded in clean studio conditions. The same file (1089-134686-0001.flac, ~3.3 s) was used across all four sub-tasks to keep results directly comparable. LibriSpeech is publicly available under CC BY 4.0 from openslr.org/12.